

Problem 1

(a) Use `datasets.make_moons` in `sklearn` package to generate a dataset with 200 samples with random noise. Visualize positive and negative instances with different colors. Split the dataset into training set and test set by 7:3

```
In [1]: # imports
import sklearn as sk
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

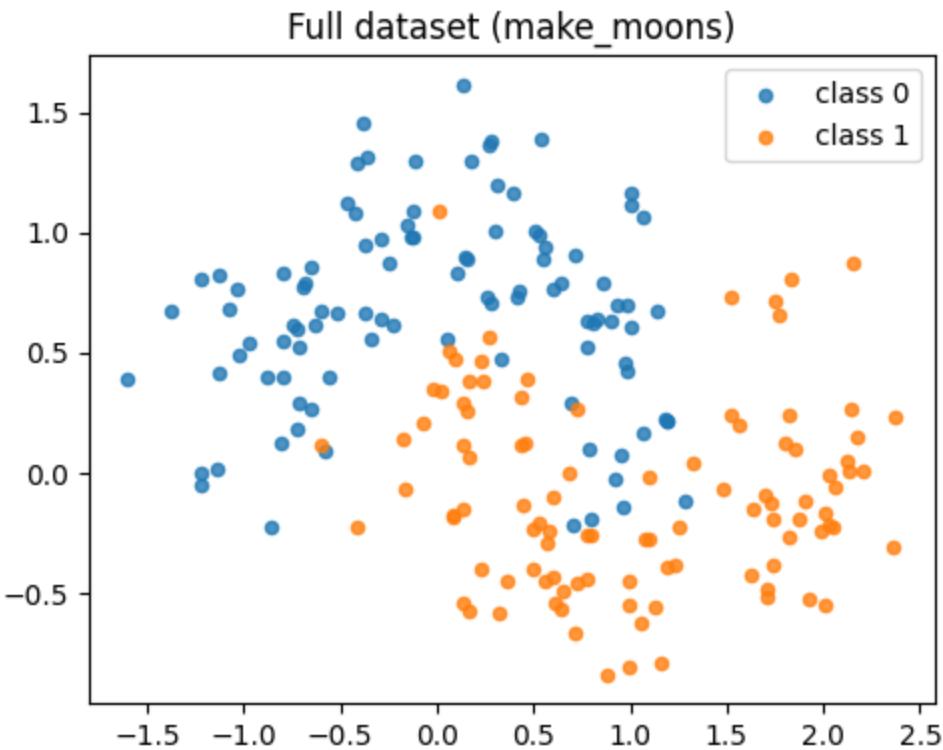
```
In [2]: # Create dataset
X, y = make_moons(n_samples=200, noise=0.25, random_state=42)
```

```
In [3]: X.shape
```

```
Out[3]: (200, 2)
```

```
In [4]: # Split dataset into training (70%) and test (30%) sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
```

```
In [5]: # Visualize dataset
plt.figure(figsize=(5,4))
plt.scatter(X[y==0,0], X[y==0,1], s=20, label="class 0", alpha=0.8)
plt.scatter(X[y==1,0], X[y==1,1], s=20, label="class 1", alpha=0.8)
plt.title("Full dataset (make_moons)")
plt.legend()
plt.tight_layout()
plt.show()
```



(b) Use SVM model in sklearn to fit the training data with “RBF” kernel. Try parameter C with different values, i.e. [0.05, 0.1, 1, 2, 5, 10, 50, 100]. Visualize the decision boundary and support vectors along with the plot in part (a). What effects does parameter C have on the model?

```
In [6]: Cs = [0.05, 0.1, 1, 2, 5, 10, 50, 100]

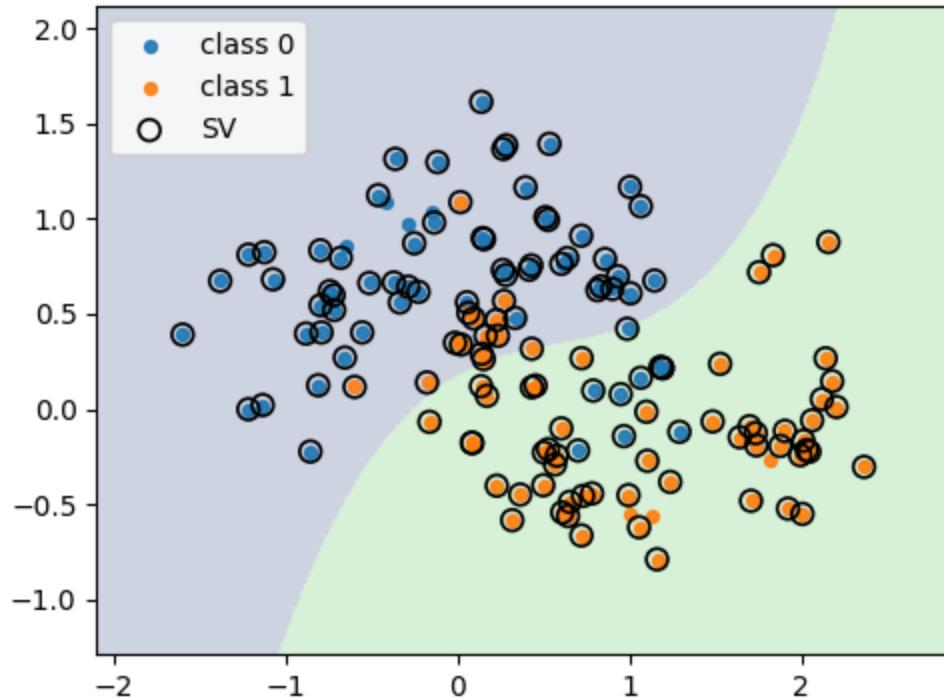
def plot_decision_boundary(clf, X, y, title):
    # Mesh
    x_min, x_max = X[:,0].min()-0.5, X[:,0].max()+0.5
    y_min, y_max = X[:,1].min()-0.5, X[:,1].max()+0.5
    xx, yy = np.meshgrid(
        np.linspace(x_min, x_max, 300),
        np.linspace(y_min, y_max, 300)
    )
    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

    # Plot
    plt.figure(figsize=(5,4))
    plt.contourf(xx, yy, Z, alpha=0.25, levels=np.arange(-0.5,2), antialiased=True)
    plt.scatter(X[y==0,0], X[y==0,1], s=18, label="class 0", alpha=0.9)
    plt.scatter(X[y==1,0], X[y==1,1], s=18, label="class 1", alpha=0.9)
    # Support vectors
    sv = clf.support_vectors_
    plt.scatter(sv[:,0], sv[:,1], s=60, facecolors='none', edgecolors='k', linewidths=2)
    plt.title(title)
    plt.legend(loc="upper left")
    plt.tight_layout()
    plt.show()

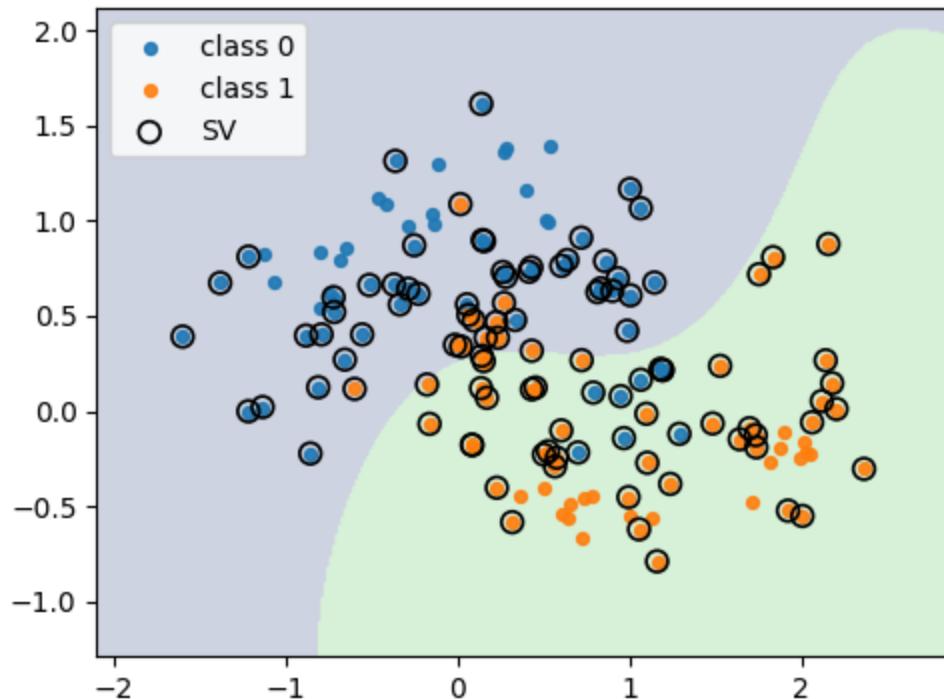
rbf_models = {}
```

```
for C in Cs:  
    clf = SVC(kernel="rbf", C=C, gamma="scale", random_state=42)  
    clf.fit(X_train, y_train)  
    rbf_models[C] = clf  
    plot_decision_boundary(clf, X_train, y_train, f"RBF SVM (C={C})")
```

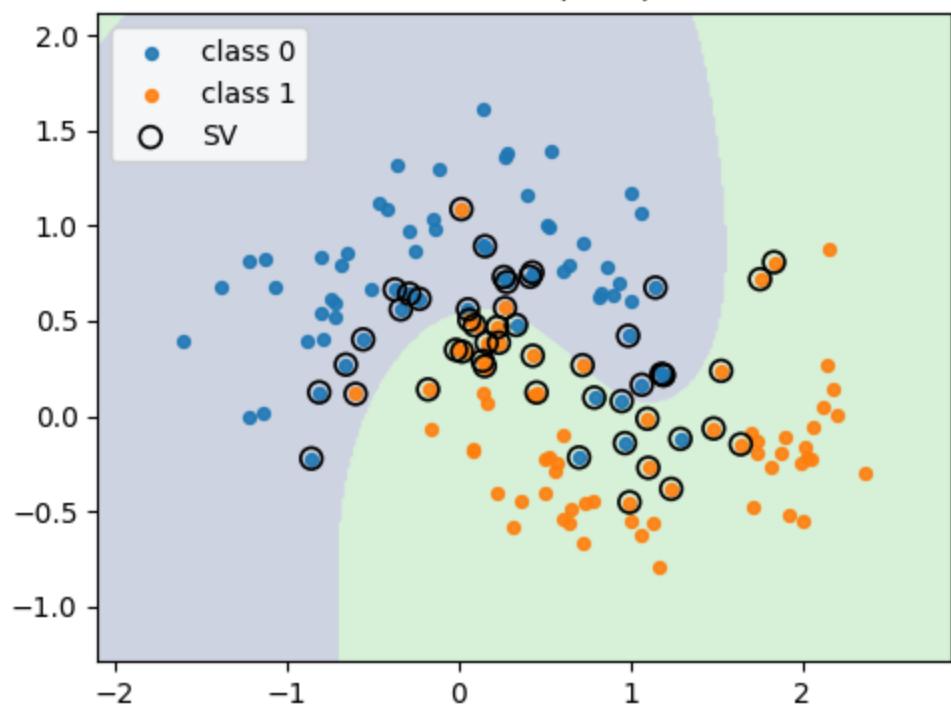
RBF SVM (C=0.05)



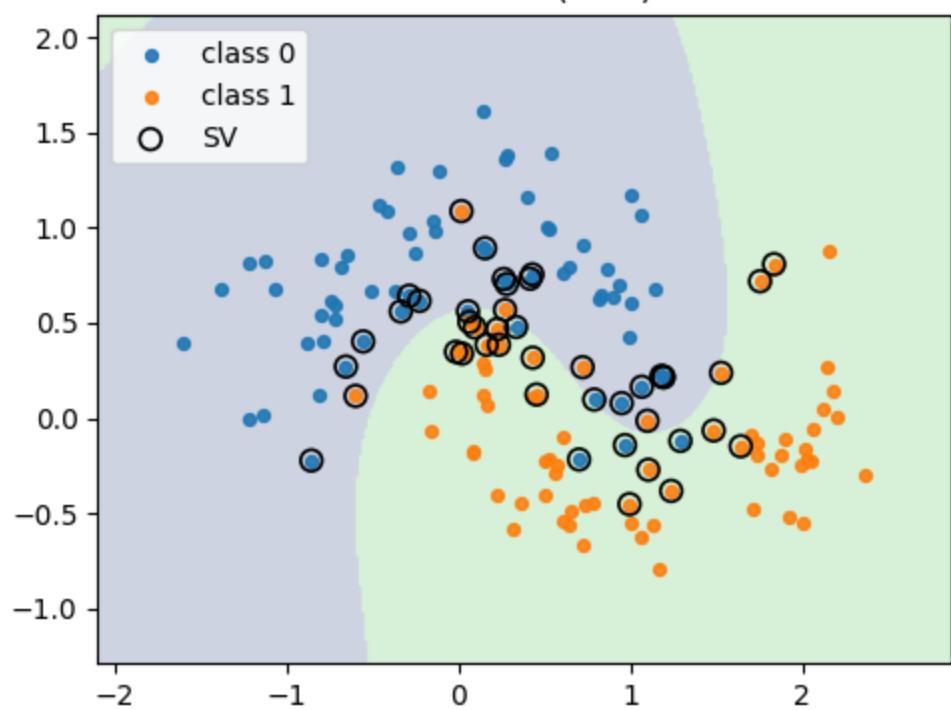
RBF SVM (C=0.1)



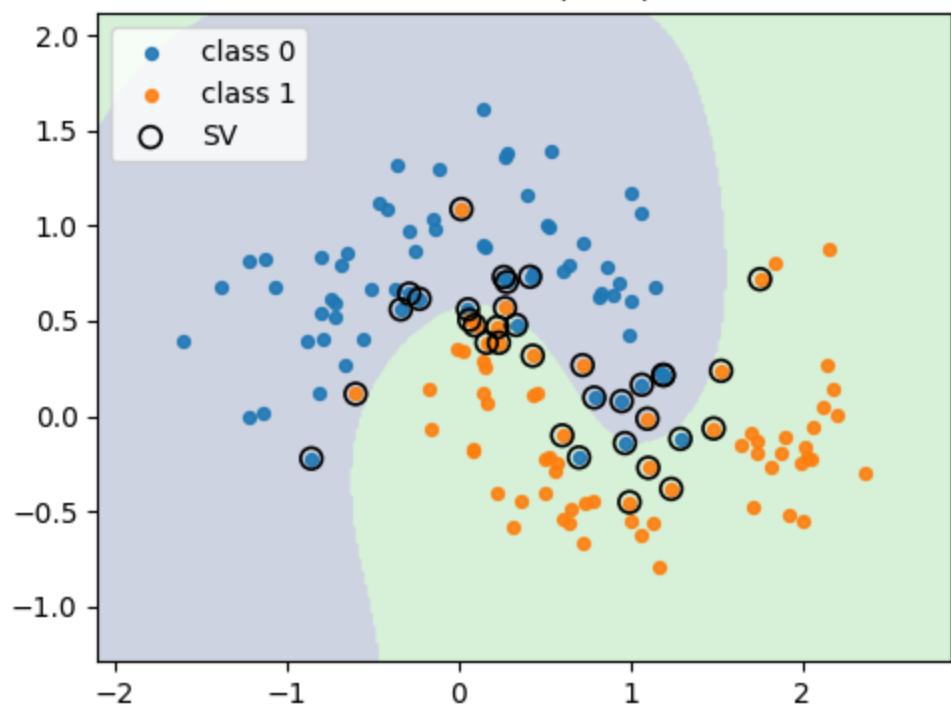
RBF SVM (C=1)



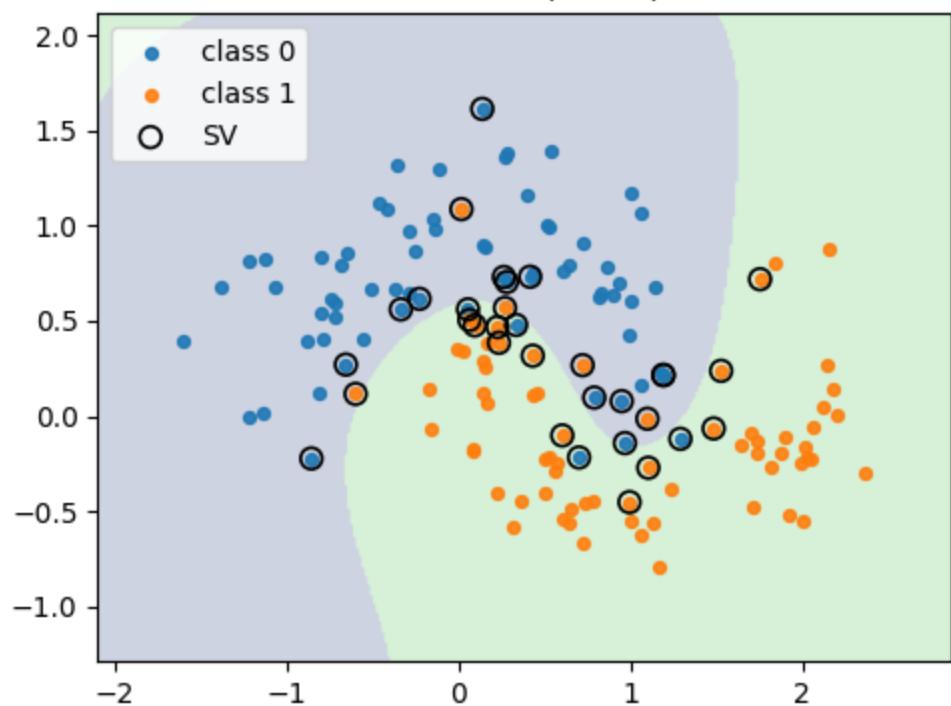
RBF SVM (C=2)



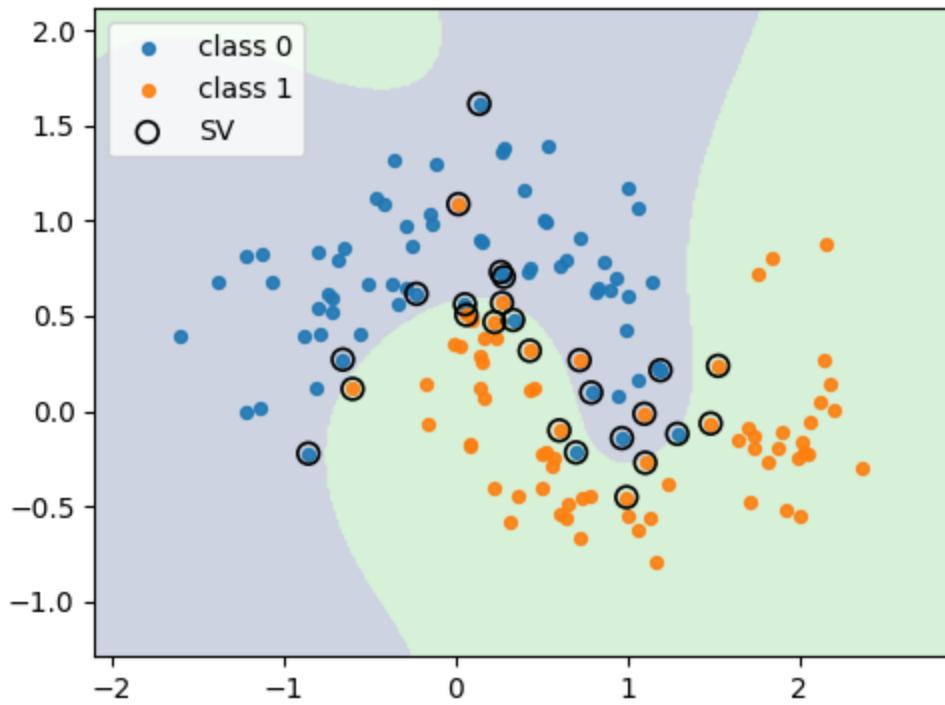
RBF SVM (C=5)



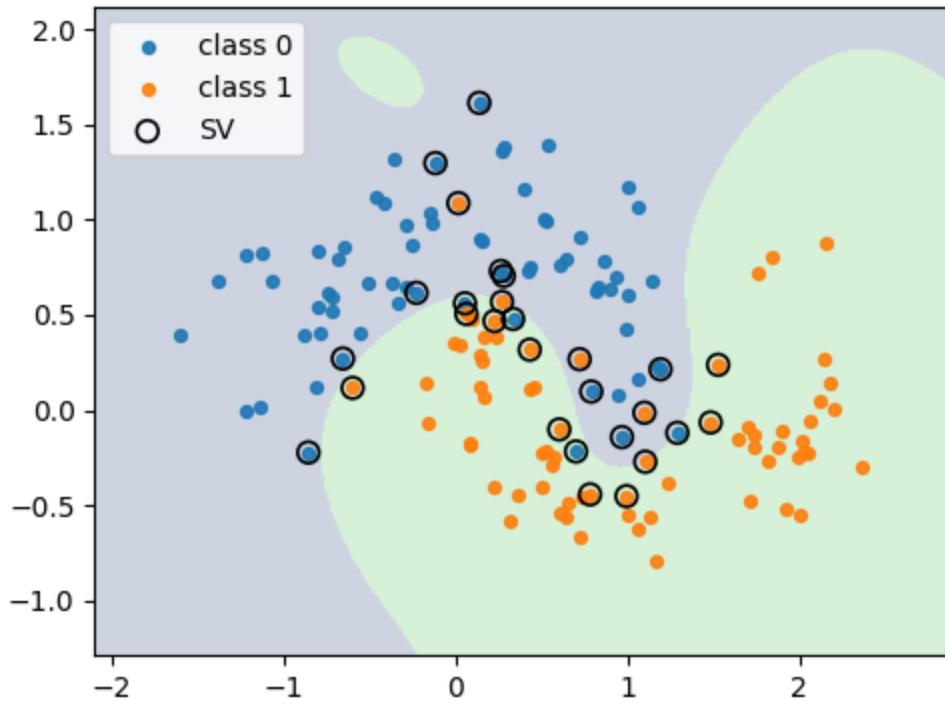
RBF SVM (C=10)



RBF SVM ($C=50$)



RBF SVM ($C=100$)



In an SVM, C controls the strength of regularization. A small C imposes stronger regularization, meaning a wider margin and more training errors. Thus, the resulting decision boundary is smoother and may underfit.

A large C weakens regularization, pushing the model to classify training points correctly with few margin violations; the boundary becomes more complex (often wiggly as seen here), which can reduce bias but increase variance and the risk of overfitting.

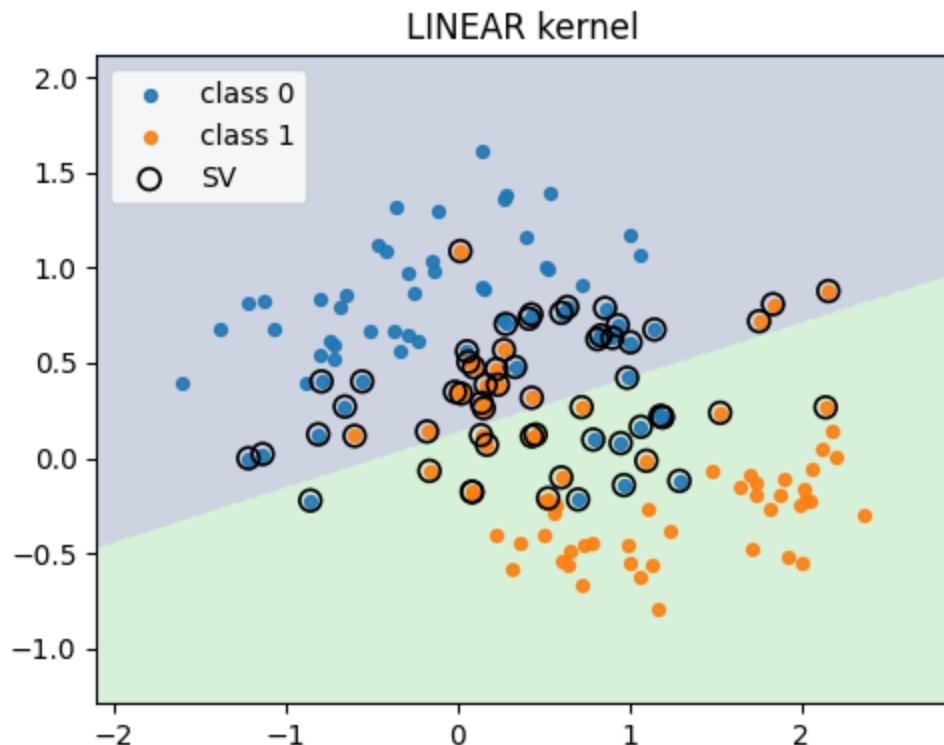
```
In [7]: # Defining kernels
kernels = [
    ("linear", {"kernel": "linear", "C": 1.0}),
    ("poly", {"kernel": "poly", "C": 1.0, "degree": 3, "gamma": "scale", "coef0": 0}),
    ("rbf", {"kernel": "rbf", "C": 5.0, "gamma": "scale"}),
    ("sigmoid", {"kernel": "sigmoid", "C": 1.0, "gamma": "scale", "coef0": 0.0}),
]
```

```
In [8]: acc = {} # accuracy of each kernel

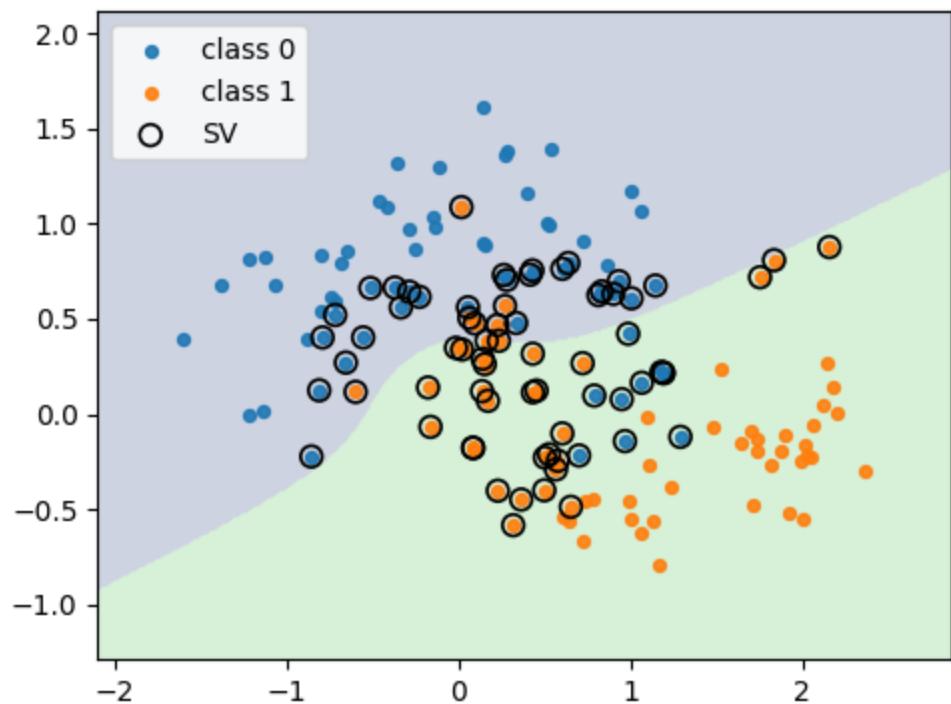
# Training and testing each kernel
for name, params in kernels:
    clf = SVC(**params, random_state=42)
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    acc[name] = accuracy_score(y_test, y_pred)
    plot_decision_boundary(clf, X_train, y_train, f"{name.upper()} kernel")

# Printing accuracy of each kernel
print("Test accuracies by kernel:")
for k, v in acc.items():
    print(f"{k:7s}: {v:.3f}")

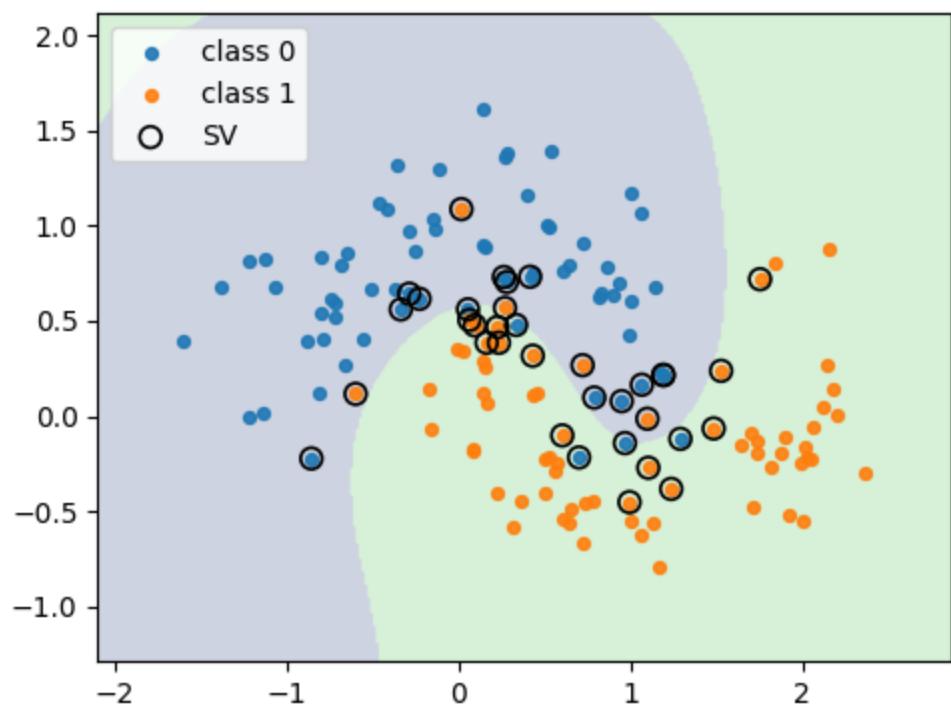
# Printing the best kernel (i.e. the one with the highest accuracy)
best_kernel = max(acc, key=acc.get)
print(f"\nBest on this split: {best_kernel.upper()} (accuracy = {acc[best_kernel]}..
```

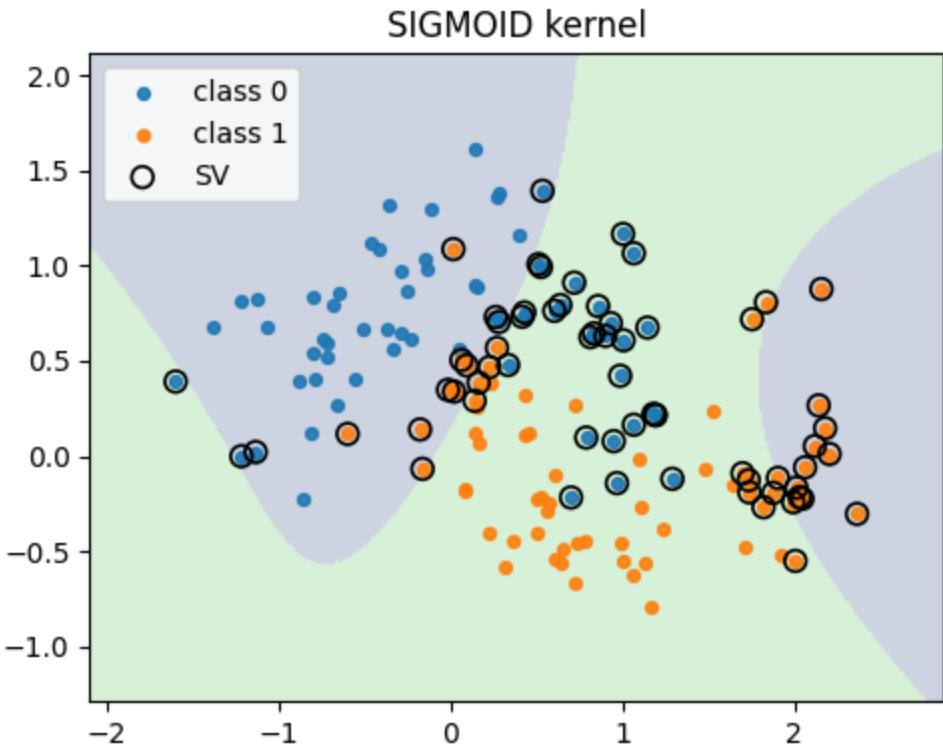


POLY kernel



RBF kernel





Test accuracies by kernel:

linear : 0.883

poly : 0.900

rbf : 0.950

sigmoid: 0.717

Best on this split: RBF (accuracy = 0.950)

Problem 2 (Artificial Neural Networks)

(1) Consider the following two activation functions:

The sigmoid function $f_1(x) = \frac{1}{1 + e^{-x}}$ and the **hyperbolic tangent** $f_2(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

Show that:

$$f'_1(x) = f_1(x)(1 - f_1(x))$$

$$f'_2(x) = 1 - f_2^2(x)$$

Claim: $f'_1(x) = f_1(x)(1 - f_1(x))$

$$\begin{aligned} f_1(x) \Rightarrow f'_1(x) &= -1 \cdot \frac{1}{(1 + e^{-x})^2} \cdot -e^{-x} = \frac{e^{-x}}{(1 + e^{-x})^2} = \frac{(1 - 1) + (1 - 1) + e^{-x}}{(1 + e^{-x})^2} \\ &= f_1(x) \left(\frac{(1 - 1) + e^{-x}}{1 + e^{-x}} \right) = f_1(x)(1 - f_1(x)) \end{aligned}$$

□

Claim: $f'_2(x) = 1 - f_2^2(x)$

$$\begin{aligned}f_2(x) \Rightarrow f'_2(x) &= \frac{(e^x + e^{-x})(e^x - (-e^{-x})) - (e^x - e^{-x})(e^x - e^{-x})}{(e^x + e^{-x})^2} = \frac{(e^x + e^{-x})^2 - (e^x - e^{-x})^2}{(e^x + e^{-x})^2} \\&= 1 - \left(\frac{e^x - e^{-x}}{e^x + e^{-x}} \right)^2 = 1 - f_2^2(x)\end{aligned}$$

□

(2) Learn how to use the MLP (multi-layer perceptron) model in sklearn package by reading corresponding sections from sklearn documentations (<https://scikit-learn.org/stable/>)

I read through the documentation

(3) Design your MLP model and apply the MLP classifier from sklearn package to the digits data set. Find the best number of hidden nodes that achieves the best classification result

```
In [9]: # imports
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import numpy as np
import pandas as pd
```

```
In [10]: # Load & split data
X, y = load_digits(return_X_y=True)

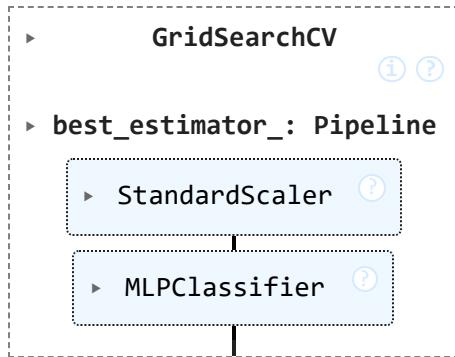
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, stratify=y, random_state=42
)
```

```
In [11]: # Pipeline: scale -> MLP (single hidden layer)
pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("mlp", MLPClassifier(
        hidden_layer_sizes=(32,),           # placeholder; tuned via grid
        activation="relu",
        solver="adam",
        alpha=1e-4,                         # L2
        learning_rate_init=1e-3,
        max_iter=1000,
        early_stopping=True,
        n_iter_no_change=20,
        random_state=42
    ))
])
```

```
    ))  
])
```

```
In [12]: # Search space (single hidden layer)  
  
candidates = [8, 16, 32, 64, 128, 192, 256, 320]  
param_grid = {"mlp_hidden_layer_sizes": [(h,) for h in candidates]}  
  
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)  
  
gs = GridSearchCV(  
    pipe,  
    param_grid=param_grid,  
    scoring="accuracy",  
    cv=cv,  
    n_jobs=-1,  
    refit=True,      # refit on full training set using the best params  
)  
  
gs.fit(X_train, y_train)
```

```
Out[12]:
```



```
In [13]: # Report best size and score  
best_size = gs.best_params_["mlp_hidden_layer_sizes"][0]  
best_cv = gs.best_score_  
  
print(f"Optimal hidden nodes (argmax CV): {best_size}")  
print(f"Best CV accuracy: {best_cv:.4f}")
```

```
Optimal hidden nodes (argmax CV): 256  
Best CV accuracy: 0.9586
```

```
In [14]: # Tie-break toward smaller networks within a tolerance  
results = pd.DataFrame(gs.cv_results_)  
results["hidden"] = results["param_mlp_hidden_layer_sizes"].apply(lambda t: t[0])  
# Find max CV score, then pick smallest hidden size within tolerance  
tol = 1e-4  
max_cv = results["mean_test_score"].max()  
candidates_within_tol = results.loc[  
    np.isclose(results["mean_test_score"], max_cv, atol=tol), "hidden"]  
]  
best_size_min = int(candidates_within_tol.min())  
  
if best_size_min != best_size:  
    print(f"Tie detected within ±{tol}; choosing smaller model: {best_size_min} nod
```

```
best_size = best_size_min

print(f"Selected optimal hidden nodes: {best_size}")
```

```
Selected optimal hidden nodes: 256
```

```
In [15]: # Test accuracy of the selected model (already refit by GridSearchCV)
test_acc = gs.best_estimator_.score(X_test, y_test)
print(f"Test accuracy: {test_acc:.4f}")
```

```
Test accuracy: 0.9815
```

```
In [16]: # Leaderboard
leaderboard = results[["hidden", "mean_test_score", "std_test_score"]].sort_values(
    ["mean_test_score", "hidden"], ascending=[False, True])
)
print("\nTop candidates:\n", leaderboard.head(5).to_string(index=False))
```

```
Top candidates:
```

hidden	mean_test_score	std_test_score
256	0.958638	0.019233
320	0.957054	0.016155
64	0.950667	0.018469
192	0.949111	0.017775
128	0.948318	0.019875

The cross-validated results indicate that a single hidden layer with 256 units performs best (mean accuracy $\approx 0.9586 \pm 0.0192$), narrowly edging out 320 units, which suggests that adding more capacity beyond 256 does not yield further gains. Accuracy generally improves as we scale from 64 \rightarrow 128 \rightarrow 192 \rightarrow 256 hidden nodes, then plateaus or slightly dips at 320, implying diminishing returns at larger sizes. The standard deviations across candidates ($\sim 0.016\text{--}0.020$) are similar, so the ranking isn't driven by fold-to-fold noise. Given both performance and efficiency, 256 hidden units is a sensible optimal choice within this search; if desired, one could probe nearby sizes (e.g., 224, 288) or repeat with multiple random seeds to confirm robustness.